TABLE I

PER-QUERY LATENCY OF THE SAFEVPR VERIFIER + CALIBRATOR OVERHEAD ON TOP OF AN EXTERNAL BACKBONE RETRIEVAL. RTX 5070 TI, FP16, PYTORCH (NO TENSORRT). T=10 FRAMES PER QUERY, P=256 PATCHES. BACKBONE RETRIEVAL IS DATASET-DATABASE-SIZE DEPENDENT AND TREATED AS BLACK-BOX.

Stage (ViT-L/14, paper default)	ms
DINOv2 query features (T=10 batched)	26.1
DINOv2 candidate features (T=10 batched)	26.1
Patch-MNN + Lowe-ratio match (T pairs, $P \times P$)	0.9
Mondrian quintile lookup + threshold compare	0.03
Total per-query overhead	**53.1**